

AI Empower Enterprise RAG Service (V5)

Deployment Manifest

This document serves as the single source of truth for the Version 5.0 architecture, encompassing all Infrastructure as Code (Terraform), application code (Python/Docker), and sequential deployment commands executed to build the fully functional, asynchronous RAG pipeline with **Storage Decoupling (V5)**.

The V5 system is deployed against the GCP Project **test-rag-backend-v4**.

1. Project Architecture and Repository Structure

1.1 V5 Architectural Summary

V5 introduces a **Dual-Write** capability to future-proof storage costs (V6 roadmap) while maintaining real-time search.

Component	V4 (Before)	V5 (Deployed State)	Purpose of Change
Ingestion Worker	Writes only to Firestore.	Writes to Firestore (Search) and GCS Parquet (Archive) .	Mitigates future cost of storing >100K vectors in Firestore.
Search Engine	Firestore Vector Search	Firestore Vector Search	Remains the primary, fast, real-time search layer.
New Dependencies	None	pandas, pyarrow	Required for efficient, columnar Parquet file writing.

1.2 Repository Structure (Complete)

ai-empower-rag-v5/

 └── .gitignore

```
└── terraform/      # Infrastructure as Code (IaC)
    ├── main.tf
    ├── variables.tf
    ├── iam.tf
    └── provider.tf
└── src/           # Cloud Run Microservices
    ├── ingestion-dispatcher/
        ├── main.py
        ├── requirements.txt
        └── Dockerfile
    ├── ingestion-worker/  # Dual-Write Worker
        ├── main.py
        ├── requirements.txt
        └── Dockerfile
    └── retrieval-api/    # REST API (Query/LLM/Memory)
        ├── main.py
        ├── requirements.txt
        └── Dockerfile
└── frontend_app/   # Streamlit Cloud App
    ├── app.py
    └── requirements.txt
```

2. Terraform Configuration (`terraform/`)

File: `terraform/provider.tf`

```
Terraform
terraform {
  required_providers {
    google = {
      source  = "hashicorp/google"
      version = ">= 5.0"
    }
  }
  provider "google" {
    project = var.project_id
    region  = var.region
  }
}
```

File: `terraform/variables.tf`

```

Terraform
variable "project_id" {
  description = "The GCP Project ID"
  type        = string
  default     = "test-rag-backend-v4"
}
variable "region" {
  description = "GCP Region"
  type        = string
  default     = "us-central1"
}
variable "bucket_name" {
  description = "Name of the V4/V5 Storage Bucket"
  type        = string
  default     = "ai-empower-rag-v4-uploads"
}

```

File: `terraform/main.tf`

```

Terraform
# 1. Storage Bucket
resource "google_storage_bucket" "rag_bucket" {
  name      = var.bucket_name
  location   = var.region
  force_destroy = false
  uniform_bucket_level_access = true
}

# 2. Pub/Sub Topic (The Async Queue)
resource "google_pubsub_topic" "ingestion_topic" {
  name = "ingestion-tasks"
}

# 3. Pub/Sub Subscription (Worker Pull/Push)
resource "google_pubsub_subscription" "ingestion_sub" {
  name = "ingestion-workers-sub"
  topic = google_pubsub_topic.ingestion_topic.name
  ack_deadline_seconds = 600
  retry_policy {
    minimum_backoff = "10s"
    maximum_backoff = "600s"
  }
}

```

```

# 4. Enable Firestore API and Create Database Instance
resource "google_project_service" "firebase" {
  service = "firebase.googleapis.com"
  disable_on_destroy = false
}
resource "google_firestore_database" "database" {
  project    = var.project_id
  name       = "(default)"
  location_id = var.region
  type       = "FIRESTORE_NATIVE"
  depends_on = [google_project_service.firebaseio]
}

# 5. Vector Index for "rag_children"
resource "google_firestore_index" "vector_index" {
  collection = "rag_children"
  database   = google_firestore_database.database.name
  fields {
    field_path = "client_id"
    order      = "ASCENDING"
  }
  fields {
    field_path = "embedding"
    vector_config {
      dimension = 768
      flat {}
    }
  }
}

```

File: `terraform/iam.tf`

Terraform

```

resource "google_service_account" "dispatcher_sa" {
  account_id  = "rag-dispatcher-sa"
  display_name = "V4 Ingestion Dispatcher"
}
resource "google_service_account" "worker_sa" {
  account_id  = "rag-worker-sa"
  display_name = "V4 Ingestion Worker"
}
# Worker Permissions (Firestore Write, Vertex AI, Storage Read)

```

```
resource "google_project_iam_member" "worker_firestore" {
  project = var.project_id
  role    = "roles/datastore.user"
  member   = "serviceAccount:${google_service_account.worker_sa.email}"
}

resource "google_project_iam_member" "worker_ai" {
  project = var.project_id
  role    = "roles/aiplatform.user"
  member   = "serviceAccount:${google_service_account.worker_sa.email}"
}
```

3. Application Services (`src/`)

3.1 Ingestion Dispatcher

File: `src/ingestion-dispatcher/requirements.txt`

Plaintext

```
functions-framework==3.*  
google-cloud-storage  
google-cloud-pubsub  
pypdf  
cloudevents
```

File: `src/ingestion-dispatcher/Dockerfile`

```
Dockerfile  
FROM python:3.11-slim  
WORKDIR /app  
COPY requirements.txt .  
RUN pip install --no-cache-dir -r requirements.txt  
COPY ..  
CMD ["functions-framework", "--target=handle_upload", "--port=8080"]
```

File: `src/ingestion-dispatcher/main.py`

```
Python  
import functions_framework  
import os  
import json  
import logging
```

```

from google.cloud import storage
from google.cloud import pubsub_v1
from pypdf import PdfReader
from io import BytesIO

PROJECT_ID = os.environ.get("GCP_PROJECT", "test-rag-backend-v4")
TOPIC_ID = "ingestion-tasks"
storage_client = storage.Client()
publisher = pubsub_v1.PublisherClient()
topic_path = publisher.topic_path(PROJECT_ID, TOPIC_ID)
logging.basicConfig(level=logging.INFO)

@functions_framework.cloud_event
def handle_upload(cloud_event):
    data = cloud_event.data
    bucket_name = data["bucket"]
    file_path = data["name"]

    if not (file_path.endswith(".pdf") or file_path.endswith(".pptx")): return

    parts = file_path.split("/")
    client_id = parts[1] if len(parts) >= 3 else "default_tenant"

    logging.info(f"Processing {file_path} for Client: {client_id}")

    bucket = storage_client.bucket(bucket_name)
    blob = bucket.blob(file_path)
    content = blob.download_as_bytes()

    try:
        if file_path.endswith(".pdf"):
            reader = PdfReader(BytesIO(content))
            total_pages = len(reader.pages)
        else:
            total_pages = 1

        logging.info(f"Found {total_pages} pages in {file_path}")

        futures = []
        for page_num in range(total_pages):
            message_json = {
                "bucket": bucket_name, "file_path": file_path, "page_num": page_num,
                "client_id": client_id, "total_pages": total_pages
            }
    
```

```
    data_str = json.dumps(message_json).encode("utf-8")
    future = publisher.publish(topic_path, data_str)
    futures.append(future)

    for f in futures: f.result()
    logging.info(f"Successfully dispatched {total_pages} tasks.")

except Exception as e:
    logging.error(f"Error dispatching {file_path}: {e}")
    raise e
```

3.2 Ingestion Worker (V5 Dual-Write)

File: `src/ingestion-worker/requirements.txt`

```
Plaintext
google-cloud-storage
google-cloud-firebase
google-cloud-pubsub
google-cloud-aiplatform
langchain==0.1.0
langchain-community==0.0.10
langchain-google-vertexai
pypdf
flask
gunicorn
pandas
pyarrow
```

File: `src/ingestion-worker/Dockerfile`

```
Dockerfile
FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY ..
CMD ["gunicorn", "--bind", ":8080", "--workers", "1", "--threads", "8", "--timeout", "0", "main:app"]
```

File: `src/ingestion-worker/main.py`

Python

```
# V5 WORKER CODE: DUAL-WRITE (Firestore for Search + Parquet for Archive)
import os
import json
import logging
import hashlib
from flask import Flask, request, jsonify
from google.cloud import storage
from google.cloud import firestore
from google.cloud.firestore_v1.vector import Vector
from langchain_google_vertexai import VertexAIEmbeddings
from langchain.text_splitter import RecursiveCharacterTextSplitter
from pypdf import PdfReader
from io import BytesIO
import base64
from google.cloud import firestore as fs_module

# V5 NEW IMPORTS
import pandas as pd
import pyarrow.parquet as pq
import pyarrow as pa
import tempfile
import time
import numpy as np

app = Flask(__name__)
logging.basicConfig(level=logging.INFO)

# --- Configuration ---
PROJECT_ID = os.environ.get("GCP_PROJECT", "test-rag-backend-v4")
ARCHIVE_BUCKET_NAME = "ai-empower-rag-v4-uploads"

db = firestore.Client(project=PROJECT_ID)
storage_client = storage.Client()
embeddings = VertexAIEmbeddings(model_name="text-embedding-004",
project=PROJECT_ID)

def get_deterministic_id(key_string):
    """Generate a hash ID to ensure Idempotency"""
    return hashlib.sha256(key_string.encode("utf-8")).hexdigest()

@app.route("/", methods=["POST"])
def process_task():
    envelope = request.get_json()
```

```

if not envelope or "message" not in envelope: return jsonify({"error": "Invalid Pub/Sub message"}), 400

pubsub_message = envelope["message"]
data_str = base64.b64decode(pubsub_message["data"]).decode("utf-8")
job = json.loads(data_str)

bucket_name = job["bucket"]
file_path = job["file_path"]
page_num = job["page_num"]
client_id = job["client_id"]

logging.info(f"Worker processing Page {page_num} of {file_path}")
parquet_data = []

try:
    bucket = storage_client.bucket(bucket_name)
    blob = bucket.blob(file_path)
    content = blob.download_as_bytes()
    reader = PdfReader(BytesIO(content))
    page = reader.pages[page_num]
    raw_text = page.extract_text()

    if not raw_text.strip(): return jsonify({"message": "Empty page text."}), 200

    parent_splitter = RecursiveCharacterTextSplitter(chunk_size=2000, chunk_overlap=200)
    parent_chunks = parent_splitter.split_text(raw_text)
    firestore_batch = db.batch()

    for p_idx, parent_text in enumerate(parent_chunks):
        parent_key = f"{file_path}|{page_num}|{p_idx}"
        parent_id = get_deterministic_id(parent_key)

        parent_ref = db.collection("rag_parents").document(parent_id)
        firestore_batch.set(parent_ref, {"client_id": client_id, "source": file_path, "page": page_num, "content": parent_text})

        child_splitter = RecursiveCharacterTextSplitter(chunk_size=400, chunk_overlap=50)
        child_chunks = child_splitter.split_text(parent_text)

        if child_chunks:
            vectors = embeddings.embed_documents(child_chunks)

            for c_idx, vector_list in enumerate(vectors):

```

```

child_key = f"{parent_key}|{c_idx}"
child_id = get_deterministic_id(child_key)
child_text = child_chunks[c_idx]

# --- 1. FIRESTORE WRITE: CHILD (for V5 search) ---
child_ref = db.collection("rag_children").document(child_id)
firestore_batch.set(child_ref, {"client_id": client_id, "parent_id": parent_id, "content": child_text, "embedding": Vector(vector_list)})

# --- 2. ACCUMULATE DATA FOR PARQUET (for V6 archive) ---
parquet_data.append({
    "id": child_id,
    "client_id": client_id,
    "parent_id": parent_id,
    "embedding": vector_list, # Store as raw Python list
    "content": child_text,
    "source": file_path,
    "page": page_num,
    "timestamp": time.time()
})

firestore_batch.commit()

# --- 3. PARQUET WRITE: ARCHIVE TO GCS ---
if parquet_data:
    df = pd.DataFrame(parquet_data)

    # Define schema for correct PyArrow list-of-float type
    schema = pa.schema([
        pa.field('id', pa.string()),
        pa.field('client_id', pa.string()),
        pa.field('parent_id', pa.string()),
        pa.field('content', pa.string()),
        pa.field('source', pa.string()),
        pa.field('page', pa.int64()),
        pa.field('timestamp', pa.float64()),
        pa.field('embedding', pa.list_(pa.float32()))
    ])

    table = pa.Table.from_pandas(df, schema=schema, preserve_index=False)

    temp_parquet_file = os.path.join(tempfile.gettempdir(),
f"page_{page_num}_{client_id}_{get_deterministic_id(file_path)}.parquet")
    pq.write_table(table, temp_parquet_file)

```

```

archive_blob_path =
f"archive/{client_id}/{os.path.basename(file_path)}/page_{page_num}.parquet"
archive_bucket = storage_client.bucket(ARCHIVE_BUCKET_NAME)
archive_blob = archive_bucket.blob(archive_blob_path)
archive_blob.upload_from_filename(temp_parquet_file)

os.remove(temp_parquet_file)
logging.info(f"V5 Archive successful for Page {page_num}: Wrote {len(parquet_data)} records to GCS Parquet.")

logging.info(f"Indexing complete for Page {page_num}.")
return jsonify({"message": f"Indexed page {page_num}"}), 200

except Exception as e:
    logging.error(f"Worker Failed: {e}")
    return jsonify({"error": str(e)}), 500

if __name__ == "__main__":
    app.run(debug=True, host="0.0.0.0", port=8080)

```

3.3 Retrieval API (V5 Memory)

File: [src/retrieval-api/requirements.txt](#)

```

Plaintext
flask
gunicorn
google-cloud-firestore
google-cloud-aiplatform
langchain==0.1.0
langchain-community==0.0.10
langchain-google-vertexai

```

File: [src/retrieval-api/Dockerfile](#)

```

Dockerfile
FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY ..
CMD ["gunicorn", "--bind", ":8080", "--workers", "1", "--threads", "8", "--timeout", "0", "main:app"]

```

File: `src/retrieval-api/main.py`

Python

```
# Implements Parent-Child fetching and conversational memory (last 4 turns)
import os
import logging
from flask import Flask, request, jsonify
from google.cloud import firestore
from google.cloud.firestore_v1.vector import Vector
from google.cloud.firestore_v1.base_vector_query import DistanceMeasure
from langchain_google_vertexai import VertexAIEmbeddings, ChatVertexAI
from langchain.prompts import ChatPromptTemplate
from google.cloud import firestore as fs_module

app = Flask(__name__)
logging.basicConfig(level=logging.INFO)

PROJECT_ID = os.environ.get("GCP_PROJECT", "test-rag-backend-v4")
CHAT_HISTORY_COLLECTION = "rag_chat_history"
MAX_HISTORY_TURNS = 4

db = firestore.Client(project=PROJECT_ID)
embeddings = VertexAIEmbeddings(model_name="text-embedding-004",
project=PROJECT_ID)
llm = ChatVertexAI(model_name="gemini-2.5-pro", project=PROJECT_ID, temperature=0.1)

@app.route("/query", methods=["POST"])
def handle_query():
    data = request.get_json()
    if not data or "query" not in data: return jsonify({"error": "Missing 'query' field"}), 400

    user_query = data["query"]
    client_id = data.get("client_id", "default_client")
    session_id = data.get("session_id", "default_session")

    try:
        # --- 1. Retrieve Chat History ---
        history_ref = db.collection(CHAT_HISTORY_COLLECTION).document(session_id)
        history_doc = history_ref.get()
        past_conversation = ""

        if history_doc.exists:
```

```

past_msgs = history_doc.to_dict().get("messages", [])[-MAX_HISTORY_TURNS:]
for msg in past_msgs:
    past_conversation += f"{msg['role']}: {msg['content']}\n"

# 2. Vector Search (Parent-Child Logic)
query_vector = embeddings.embed_query(user_query)
collection = db.collection("rag_children")
vector_ref = Vector(query_vector)

results = collection.where(filter=firestore.FieldFilter("client_id", "==", client_id))\
    .find_nearest(
        vector_field="embedding",
        query_vector=vector_ref,
        distance_measure=DistanceMeasure.COSINE,
        limit=7,
        distance_result_field="distance"
    ).get()

context_text = "No relevant documents found."
if results:
    parent_ids = set(doc.to_dict()["parent_id"] for doc in results)
    parent_refs = [db.collection("rag_parents").document(pid) for pid in parent_ids]
    parent_docs = db.get_all(parent_refs)

    context_text = ""
    for p_doc in parent_docs:
        if p_doc.exists:
            chunk = p_doc.to_dict()
            context_text += f"\n[Source: {chunk['source']}, Page: {chunk['page']}]\n{chunk['content']}\n"

# 3. Generate Answer
prompt = ChatPromptTemplate.from_template("""
    You are an expert medical AI assistant. Answer the question strictly based on the provided
    context.

    Use the Chat History to understand follow-up questions or resolve ambiguous pronouns
    (e.g., "the two").
""")


```

Chat History:
 {history}

Context:
 {context}

```

Question:
{question}

Answer:
""")

chain = prompt | llm
response = chain.invoke({
    "context": context_text,
    "history": past_conversation if past_conversation else "No previous conversation history.",
    "question": user_query
})

final_response = {"answer": response.content, "context_used": context_text}

# --- 4. Save State (Update History) ---
new_messages = [
    {"role": "user", "content": user_query},
    {"role": "assistant", "content": response.content}
]
history_ref.set({"messages": fs_module.ArrayUnion(new_messages)}, merge=True)

return jsonify(final_response), 200

except Exception as e:
    logging.error(f"Retrieval Error: {e}")
    return jsonify({"error": str(e)}), 500

if __name__ == "__main__":
    app.run(debug=True, host="0.0.0.0", port=8080)

```

4. Frontend ([frontend_app/](#))

File: [frontend_app/requirements.txt](#)

```

Plaintext
streamlit
requests
google-cloud-storage
google-auth

```

File: frontend_app/app.py

Python

```
import streamlit as st
import requests
import json
import uuid
from google.cloud import storage
from google.oauth2 import service_account
from google.auth import exceptions as auth_exceptions

# --- V4 CONFIGURATION ---
PROJECT_ID = "test-rag-backend-v4"
BUCKET_NAME = "ai-empower-rag-v4-uploads"
API_URL = "https://rag-retrieval-v4-873142271416.us-central1.run.app/query"

st.set_page_config(page_title="AI Empower RAG V4", layout="wide")
st.title("Enterprise RAG V4 (Async Parent-Child)")
st.subheader("Client Project: " + PROJECT_ID)

# --- State Management ---
if "session_id" not in st.session_state:
    st.session_state.session_id = str(uuid.uuid4())
if "messages" not in st.session_state:
    st.session_state.messages = []
if "client_id" not in st.session_state:
    st.session_state.client_id = "test_client"

@st.cache_resource(ttl=3600)
def get_storage_client():
    if "gcp_service_account" not in st.secrets:
        st.warning("GCP Service Account Secret is required for upload and not configured.")
        return None

    try:
        key_dict = dict(st.secrets["gcp_service_account"])
        creds = service_account.Credentials.from_service_account_info(key_dict)
        return storage.Client(credentials=creds, project=PROJECT_ID)
    except Exception as e:
        st.error(f"Credential setup failed. Check st.secrets['gcp_service_account']: {e}")
        return None

with st.sidebar:
    st.header("Tenant & Data Management")
```

```

st.session_state.client_id = st.text_input("Client ID (Tenant Key)",
                                         value=st.session_state.client_id)

st.header("Self-Service Upload")
uploaded_file = st.file_uploader("PDF/PPTX Document", type=['pdf', 'pptx'])

if uploaded_file and st.button("Upload & Ingest"):
    client = get_storage_client()
    if client:
        with st.spinner(f"Uploading to {st.session_state.client_id}..."):
            try:
                bucket = client.bucket(BUCKET_NAME)
                blob_path = f"uploads/{st.session_state.client_id}/{uploaded_file.name}"
                blob = bucket.blob(blob_path)
                blob.upload_from_file(uploaded_file, rewind=True)

                st.success("Upload Complete! Indexing started (V4 Async Pipeline).")
                st.info("The document will be searchable in a few minutes.")

            except Exception as e:
                st.error(f"Upload Failed. Check permissions on the bucket: {e}")

for msg in st.session_state.messages:
    st.chat_message(msg["role"]).write(msg["content"])

if prompt := st.chat_input("Ask a question about your documents..."):
    st.session_state.messages.append({"role": "user", "content": prompt})
    st.chat_message("user").write(prompt)

    with st.chat_message("assistant"):
        with st.spinner(f"Consulting {st.session_state.client_id} knowledge base..."):
            try:
                response = requests.post(API_URL, json={
                    "query": prompt,
                    "client_id": st.session_state.client_id,
                    "session_id": st.session_state.session_id
                })

                if response.status_code != 200:
                    error_data = response.json()
                    st.error(f"API Error ({response.status_code}): {error_data.get('error', 'Unknown Error')}")
                    answer = "Error processing request."
            except requests.exceptions.RequestException as e:
                st.error(f"Request failed: {e}")
                answer = "Error processing request."
else:
    st.chat_message("assistant").write("Hello! How can I assist you today?")

```

```
data = response.json()
answer = data.get("answer", "Error retrieving answer.")
st.markdown(answer)

st.session_state.messages.append({"role": "assistant", "content": answer})

except Exception as e:
    st.error(f"Connection or runtime error: {e}")
```